

---

# An Entropy Maximizing Geohash for Distributed Spatiotemporal Database Indexing

---

**Taylor B. Arnold**  
AT&T Labs Research  
33 Thomas Street  
New York, NY 10007  
`taylor@research.att.com`

## Abstract

We present a modification of the standard geohash algorithm based on maximum entropy encoding in which the data volume is approximately constant for a given hash prefix length. Distributed spatiotemporal databases, which typically require interleaving spatial and temporal elements into a single key, reap large benefits from a balanced geohash by creating a consistent balance between spatial and temporal precision even across areas of varying data density. This property also useful for indexing purely spatial datasets, where load distribution of large range scans is an important aspect of query performance. We apply our algorithm to data generated proportional to population as given by census block population counts provided from the US Census Bureau. An efficient implementation for calculating an arbitrary balanced geohash is also provided.

## 1 Introduction

There has been a rapid increase in petabyte-scale data sources across a range of industry, government, and academic applications. Many of these sources include data tagged with both geospatial and temporal attributes. Website HTTP request logs, along with geocoded IP addresses, are a common example where even modestly sized sites can quickly produce large-scale datasets. Inexpensive embedded GPS devices are another common source, and are typically found in everything from automobiles to cellphones. Sensor data from Radio-frequency identification (RFID) devices, which have been readily adapted to monitor complex supply-chain management operations, amongst other applications, also generate large datasets with both location and time-based features.

General purpose distributed filesystems such as the Google Filesystem [9] and the Hadoop Filesystem [15], provide adequate support for the raw storage of geospatial data sources. These alone however only provide minimal support for querying file chunks. Analysis, aggregations, and filtering often require processing through a large subset of the raw data. Database functionality is provided by additional software, often modelled off of Google's BigTable [3] design. Data are stored lexicographically by a primary key; queries take the form of a scan over a range of contiguous primary keys along with optional, additional query filters. Popular open source implementations of this form include Amazon's DynamoDB [5], Apache Accumulo [8], Apache HBase [17], and Apache Cassandra [12].

For geospatial data sources which are typically queried on an entity-by-entity basis, the standard distributed database applications are straightforward to implement. A uniquely identifying serial number for each entity can be used as the primary key, with spatial and temporal filtering accomplished with a combination of in-database filters as a single entity will typically constitute only a negligible proportion of the overall data volume. This application model for example would satisfy the needs of a user-facing supply chain system in which end-users search for the current location

of an expected delivery. Unfortunately, more complex spatial queries do not fit neatly into the distributed database design. A single primary key is natively one-dimensional, whereas geospatial data is typically two-dimensional and, with the addition of a temporal component, often no less than three. A common method for alleviating the dimensionality problem of having a single primary index is through the use of a space-filling curves such as Peano, Gosper, and Dragon curves. A geohash is a particular algorithm which uses a space-filling curve for encoding latitude and longitude tuples as a single string. It has had fairly wide adoption for this purpose; for example, it has been included as a core functionality in the popular distributed document store MongoDB [4].

The geohash concept adapts well for purely geospatial distributed datasets, allowing for querying spatial regions of vastly different scales through the use of a single index; larger regions specify a longer range of geohashes whereas smaller regions specify a shorter range. A temporal component can be incorporated into the space-filling curve underlying a standard geohash, however this requires an unavoidable, fixed database design decision as to the appropriate scaling factor between space and time. For instance: should a square kilometre bucket in space correspond to a month, day, or hour bucket in time? The decision has important performance implications as spatial queries take the form of a collection of range queries of primary keys. If the buckets are not appropriately sized queries will either include many false positives (data returned outside of the area of interest) or require stitching together a large number of small scans. Either of these can cause significant performance deterioration, due to additional I/O time in the case of false positives and additional overhead in running a large number of small jobs.

In many cases the question of how to balance space and time into a single key is complicated by the fact that data are distributed very unevenly over the spatial dimensions. Sources which are generated roughly proportional to population density will produce 4 orders of magnitude or more in urban environments compared to rural ones. In these common cases, it is often more natural to scale the spatial buckets to be proportional to the data volume. In this case, we instead only need to decide: should a bucket with 1 million data points correspond to month, day, or hour bucket in time? This question will typically have a more natural answer for a given database system than the previous one, particularly when the query results are used as inputs to complex algorithm such as recommender systems or anomaly detection routines.

In order to implement a data balanced alternative to the standard space filling curves, we present a novel entropy maximizing geohash. Our method is necessarily model based and learned from a small subset of spatial data. It is robust both to random noise and model drift. It can be implemented with only a modest increase in computational complexity and is generic enough to be implemented as-is in conjunction with several recently proposed geohash based schemes for storing large geospatial-temporal datasets. As a side effect of the entropy based encoding it has also reduced the size of the raw database files; this has been observed even after applying aggressive compression techniques.

The remainder of this article is organized as follows: Section 2 provides a brief literature review of geospatial database techniques with a particular focus on how these can be used in conjunction with our methods. In Section 3 we give a thorough mathematical description of the standard geohash algorithm and present the formulation for entropy maximizing geohash. Theorem 1 establishes an upper bound on how well the maximized geohash can be learnt from a sample of spatial data. Section 4 applies the entropic geohash to data based on data from the US Census Bureau, and empirically demonstrates how robust the method is to model drift. Section 5 gives results from an implementation of the geohash in HBase using synthetic data generated from the Census bureau and Section 6 provides avenues for future extensions.

## 2 Related Work

There is a long history of database work on data structures for optimizing geospatial and geospatial-temporal queries such as  $k$ -nearest neighbors and spatial joins. Common examples include R-trees [10],  $R^*$ -trees [2], and quadrees [14]. These algorithms often offer superior performance to simple space filling curves by replicating the true dimensionality of the problem by an abstraction of a tree as linked lists. Unfortunately linked lists are do not fit into the primary key design scheme of the distributed databases mentioned in the previous section making these more sophisticated algorithms out of reach for productionalized systems.

Some research has been done towards producing a distributed database which would be able natively implement the geospatial temporal database structures which have been successful on single machine scale applications. These include VegaGiStore [20], VegaCache [18] and VegaCI [19], [21]. However at this time these remain merely theoretical as no distribution of any of these has even been publicly distributed, let alone developed to the maturity needed for a production system.

The simpler approach of adapting geohashes by interleavings spatial and temporal dimensions has also been explored. Given that these can be quickly implemented on top of existing industry standard databases this is the approach we have choosen to augment with our work. Fox et al. developed a generic approach for blending these two dimensions, along with methods for extending to lines and polygons and for producing fast queries in Apache Accumulo [7]. The HGrid data model is a similar application which takes advantage of the specific secondary indicies and related optimizations present in HBase [11]. These both rely on specifying the relative resolution of the spatial and temporal dimensions, and would would be adaptable with no major changes to use our entropy maximizing geohash.

Finally, we note that the desire for a data balanced design which motivated our work is a natural consequence of effcient query design, and has driven much of the advanced tree-based work. Red-black trees [1] and the Day-Stout-Warren algorithm for binary trees [16] are two important examples. In this context, the entropy maximizing geohash can be thought of as the curve filling analogue to balanced trees.

### 3 Entropy Balanced Geohash

#### 3.1 A Formulation of the Standard Geohash Encoding

As mentioned, a geohash is a simple scheme for mapping two-dimensional coordinates into a hierarchical, one-dimensional encoding. It is explained in several other sources, but we re-construct it here in a format which will be most conducive to generalizations. The first step is to map latitude and longitude coordinates in a standard unit square; this is done by the following linear mapping:

$$x = \frac{\text{lon} + 180}{360} \quad (1)$$

$$y = \frac{\text{lat} + 90}{180} \quad (2)$$

This choice is by convention, and any other method for mapping coordinantes into the unit square is equally viable. The geohash formulation actually works on any continuous two-dimensional data, and could be used as-is to encode other two-dimensional datasets.

The  $x$  and  $y$  coordinates need to be expressed in as a binary decimals. Formally, we define the  $x_i \in \{0, 1\}$  and  $y_i \in \{0, 1\}$  (with the restriction that neither is allowed to have an infinite trailing tail with all '1's, for uniqueness) such that

$$x = \sum_{i=1}^{\infty} \frac{x_i}{2^i}, \quad (3)$$

$$y = \sum_{i=1}^{\infty} \frac{y_i}{2^i}. \quad (4)$$

A geohash representation of  $(x, y)$  is constructed by interleaving these binary digits. The  $q$ -bit geohash  $g_q(\cdot, \cdot)$  can symbolically be defined as

$$g_q(x, y) := \sum_{i=1}^{\lceil q/2 \rceil} \frac{x_i}{2^{2i-1}} + \sum_{i=1}^{\lfloor q/2 \rfloor} \frac{y_i}{2^{2i}}. \quad (5)$$

It is fairly easy to show that the geohash function is monotone increasing in  $q$ , with the growth strictly bounded by  $2^{-q}$ , so that

$$0 \leq g_{q+m}(x, y) - g_q(x, y) < \frac{1}{2^q} \quad (6)$$

For all  $m$  greater than zero.

### 3.2 Entropy

A geohash is typically used as an index in the storing and querying of large spatial processes. A simple theoretical model for a stream of spatial data can be constructed by assuming that each observation is an independent identically distributed random variable  $\mathfrak{F}$  from some distribution over space. Borrowing a concept from information theory, we can define the entropy of a geohash over a given spatial distribution by the equation

$$H(g_q) := -1 \sum_{v \in \mathcal{R}(g_q)} \mathbb{P}[g_q(\mathfrak{F}) = v] \cdot \log_2 \{\mathbb{P}[g_q(\mathfrak{F}) = v]\} \quad (7)$$

Where  $\mathcal{R}(g_q)$  is the range of the  $q$ -bit geohash. It is a standard result that the entropy of a discrete distribution is maximized by the uniform distribution. Therefore we can use this as a proxy for how balanced a geohash is for a given distribution of spatial data.

### 3.3 The Generalized Geohash

As the  $q$ -bit geohash function is bounded and monotonic in  $q$ , we can define the infinite precision geohash, which we denote as simply  $g(\cdot, \cdot)$ , to be the limit

$$\lim_{q \rightarrow \infty} g_q(x, y) := g(x, y). \quad (8)$$

With this continuous format, one can see that if we compose  $g$  with an appropriate new function  $h$ , the composition  $h \circ g(x, y)$  can be thought of as a rescaled version of the traditional geohash. To be precise, we would like a function  $h$  to have the following properties:

$$h : [0, 1] \rightarrow [0, 1], \quad (9)$$

$$h(0) = 0, \quad (10)$$

$$h(1) = 1, \quad (11)$$

$$x < y \iff h(x) < h(y). \quad (12)$$

Note that Equation 12 implies that  $h$  is also continuous. From here, we can define the analogue to a  $q$ -bit geohash by truncating the binary representation of  $h(z) = w$ ,

$$h(z) = \sum_{i=1}^{\infty} \frac{w_i}{2^i} \quad (13)$$

To the its first  $q$ -bits

$$h_q(z) := \sum_{i=1}^q \frac{w_i}{2^i}. \quad (14)$$

In the remainder of this document, we refer the  $h_q \circ g(x, y)$  as a generalized geohash.

### 3.4 The Empirical Entropic Geohash

We have introduced the concept of a generalized geohash in order to construct a spatial encoding scheme which better optimizes the entropy as defined in Equation 7. Assume  $\{z_i\}_{i=0}^N$  is a set of independent sample from realizations of the random variable  $\mathfrak{F}$ . The empirical cumulative distribution function  $G$  of the standard geohash function  $g(\cdot, \cdot)$  is given by

$$G(t) := \frac{1}{N} \cdot \sum_{i=0}^N 1_{g(z_i) \leq t}, \quad t \in [0, 1]. \quad (15)$$

From Equation 15 we can define the balanced entropy maximizing geohash function  $b$  (balanced), assuming that every point  $z_i$  has a geohash  $g(z_i)$  strictly between 0 and 1, to be

$$b^q(t) := \begin{cases} 0 & \text{if } t = 0 \\ 1 & \text{if } t = 1 \\ \frac{N}{N+2} \cdot G^{-1}(t) & \text{if } \exists i \in \mathbb{Z} \text{ s.t. } t = i/2^q \\ \text{linear interpolation of the above points} & \text{else} \end{cases} \quad (16)$$

The balanced geohash is essentially the inverse of the empirical function  $G$ , with some minor variations to satisfy Equations 9-12. If the points  $\{z_i\}$  are unique, and  $N$  is sufficiently large, the  $q$ -bit analogue  $b_q^q$  of Equation 16 will have an entropy  $H(b_q^q)$  of approximately equal to  $q$ .

More formally, we can prove the following bound on the entropy of the balanced geohash:

**Theorem 1.** *Let  $b_q^q$  be the entropy balanced geohash estimated from a sample of  $N$  unique data points. Then, with probability at least  $1 - 2e^{-0.49 \cdot 2^{-2q} N \cdot [1-A]^2}$  the entropy  $H(b_q^q)$  is bounded by the following simultaneously for all values  $A \in [0, 1]$ :*

$$H(b_q^q) \geq q \cdot \frac{N}{N+2} \cdot A \quad (17)$$

*Proof.* Let  $F(\cdot)$  be the true cumulative distribution function of the variable  $\mathfrak{F}$ , and  $F_N(\cdot)$  be the empirical cumulative distribution function from a sample of  $N$  independent observations. Setting  $\epsilon = [1 - A] \cdot \frac{N}{N+2} \cdot 2^{-(q+1)}$ , the Dvoretzky-Kiefer-Wolfowitz inequality states [6] that the follow holds for all values  $A \in [0, 1]$  with probability  $1 - e^{-2N\epsilon}$ :

$$|F(x) - F_n(x)| \leq [1 - A] \cdot \frac{N}{N+2} \cdot 2^{-(q+1)} \quad (18)$$

Therefore, the empirical entropy should be bounded as follows:

$$H(b_q^q) \geq -1 \cdot \sum_{i=0}^{2^q-1} [F(i/2^q) - F((i+1)/2^q)] \times \log_2 [F(i/2^q) - F((i+1)/2^q)] \quad (19)$$

$$\geq -1 \cdot \sum_{i=0}^{2^q-1} [F_n(i/2^q) - F_n((i+1)/2^q) - 2\epsilon] \times \log_2 [F_n(i/2^q) - F_n((i+1)/2^q) - 2\epsilon] \quad (20)$$

$$= -2^q \cdot \left[ \frac{N}{N+2} \cdot \frac{1}{2^q} - 2\epsilon \right] \times \log_2 \left[ \frac{N}{N+2} \cdot \frac{1}{2^q} - 2\epsilon \right] \quad (21)$$

$$= -1 \cdot \frac{N}{N+2} A \times \log_2 \left[ \frac{N}{N+2} \cdot \frac{1}{2^q} \cdot A \right] \quad (22)$$

$$\geq q \cdot \frac{N}{N+2} A \quad (23)$$

Which, plugging in the appropriate  $\epsilon$  into the probability bound, yields the desired result.  $\square$

Plugging in  $A = 2/3$ , the bound in Theorem 1 holds with probability greater than 0.99 for a 5-bit balanced geohash when  $N$  is at least  $1e5$  and for a 10-bit geohash when  $N$  is at least  $1e8$ . We will see in our empirical examples that the rate of convergence of the entropy to  $q$  is typically much faster.

## 4 Census Data Example

### 4.1 Population based

## 5 Distributed Application

## 6 Future Work

## References

- [1] Rudolf Bayer. Symmetric binary b-trees: Data structure and maintenance algorithms. *Acta informatica*, 1(4):290–306, 1972.
- [2] Norbert Beckmann, Hans-Peter Kriegel, Ralf Schneider, and Bernhard Seeger. *The R\*-tree: an efficient and robust access method for points and rectangles*, volume 19. ACM, 1990.



Figure 1: Example entropy balanced geohash buckets (those which have a common prefix) of 3,4,5, and 6-bits learned from population data in the state of Ohio. Numbers are sequentially ordered from the smallest to the largest geohash buckets, and the buckets were truncated to the state boundaries to increase readability even though they theoretically extend outside as well.

- [3] Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C Hsieh, Deborah A Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E Gruber. Bigtable: A distributed storage system for structured data. *ACM Transactions on Computer Systems (TOCS)*, 26(2):4, 2008.
- [4] Kristina Chodorow. *MongoDB: the definitive guide*. ” O’Reilly Media, Inc.”, 2013.
- [5] Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall, and Werner Vogels. Dynamo: amazon’s highly available key-value store. In *ACM SIGOPS Operating Systems Review*, volume 41:6, pages 205–220. ACM, 2007.
- [6] Aryeh Dvoretzky, Jack Kiefer, and Jacob Wolfowitz. Asymptotic minimax character of the sample distribution function and of the classical multinomial estimator. *The Annals of Mathematical Statistics*, pages 642–669, 1956.
- [7] Anthony Fox, Chris Eichelberger, James Hughes, and Skylar Lyon. Spatio-temporal indexing in non-relational distributed databases. In *Big Data, 2013 IEEE International Conference on*, pages 291–299. IEEE, 2013.
- [8] Adam Fuchs. Accumulo—extensions to googles bigtable design. Technical report, Technical report, National Security Agency, 2012.
- [9] Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. The google file system. In *ACM SIGOPS Operating Systems Review*, volume 37, pages 29–43. ACM, 2003.
- [10] Antonin Guttman. *R-trees: a dynamic index structure for spatial searching*, volume 14. ACM, 1984.
- [11] Dan Han and Eleni Stroulia. Hgrid: A data model for large geospatial data sets in hbase. In *Proceedings of the 2013 IEEE Sixth International Conference on Cloud Computing*, pages 910–917. IEEE Computer Society, 2013.
- [12] Avinash Lakshman and Prashant Malik. Cassandra: a decentralized structured storage system. *ACM SIGOPS Operating Systems Review*, 44(2):35–40, 2010.
- [13] Youzhong Ma, Yu Zhang, and Xiaofeng Meng. St-hbase: a scalable data management system for massive geo-tagged objects. In *Web-Age Information Management*, pages 155–166. Springer, 2013.
- [14] Hanan Samet and Robert E Webber. Storing a collection of polygons using quadrees. *ACM Transactions on Graphics (TOG)*, 4(3):182–222, 1985.
- [15] Konstantin Shvachko, Hairong Kuang, Sanjay Radia, and Robert Chansler. The hadoop distributed file system. In *Mass Storage Systems and Technologies (MSST), 2010 IEEE 26th Symposium on*, pages 1–10. IEEE, 2010.
- [16] Quentin F. Stout and Bette L. Warren. Tree rebalancing in optimal time and space. *Communications of the ACM*, 29(9):902–908, 1986.
- [17] Ronald C Taylor. An overview of the hadoop/mapreduce/hbase framework and its current applications in bioinformatics. *BMC bioinformatics*, 11(Suppl 12):S1, 2010.
- [18] Yunqin Zhong, Jinyun Fang, and Xiaofang Zhao. Vegacache: Efficient and progressive spatio-temporal data caching scheme for online geospatial applications. In *Geoinformatics (GEOINFORMATICS), 2013 21st International Conference on*, pages 1–7. IEEE, 2013.
- [19] Yunqin Zhong, Jizhong Han, Tieying Zhang, and Jinyun Fang. A distributed geospatial data storage and processing framework for large-scale webgis. In *Geoinformatics (GEOINFORMATICS), 2012 20th International Conference on*, pages 1–7. IEEE, 2012.
- [20] Yunqin Zhong, Jizhong Han, Tieying Zhang, Zhenhua Li, Jinyun Fang, and Guihai Chen. Towards parallel spatial query processing for big spatial data. In *Parallel and Distributed Processing Symposium Workshops & PhD Forum (IPDPSW), 2012 IEEE 26th International*, pages 2085–2094. IEEE, 2012.
- [21] Yunqin Zhong, Xiaomin Zhu, and Jinyun Fang. Elastic and effective spatio-temporal query processing scheme on hadoop. In *Proceedings of the 1st ACM SIGSPATIAL International Workshop on Analytics for Big Geospatial Data*, pages 33–42. ACM, 2012.

	Geohash						Balanced Hash								
	Population			Rentals			Population			Rentals					
	$H(g_{10})$	$H(g_{20})$	$H(g_{30})$	$H(g_{10})$	$H(g_{20})$	$H(g_{30})$	$H(h_{10}^{10})$	$H(h_{20}^{10})$	$H(h_{30}^{10})$	$H(h_1^{10})$	$H(h_5^{10})$	$H(h_{10}^{10})$	$H(h_{20}^{10})$	$H(h_{30}^{10})$	
DC	0.0	1.2	7.6	0.0	1.2	7.2	9.8	11.3	11.3	1.0	4.9	9.4	10.5	10.5	
DE	1.0	3.6	10.8	0.9	3.2	9.5	9.9	12.9	12.9	1.0	4.9	9.3	11.5	11.5	
HI	0.8	3.4	10.0	0.8	3.3	9.3	9.9	12.2	12.3	1.0	4.9	9.6	11.6	11.7	
AK	1.9	5.4	11.1	1.9	4.9	10.4	9.9	12.3	12.4	1.0	4.9	9.6	11.6	11.8	
ND	0.8	5.7	11.6	0.7	4.6	9.8	10.0	13.7	14.1	1.0	4.9	9.4	11.9	12.0	
VT	0.0	5.3	12.0	0.0	5.0	10.6	9.9	12.9	12.9	1.0	4.9	9.6	12.0	12.0	
RI	0.0	2.6	10.4	0.0	2.4	9.3	10.0	13.3	13.3	1.0	4.9	9.6	12.4	12.4	
NH	0.0	4.7	12.4	0.0	4.5	10.7	10.0	13.7	13.7	1.0	4.9	9.5	12.5	12.5	
WY	0.0	5.7	11.2	0.0	5.4	10.2	10.0	13.4	13.6	1.0	4.9	9.6	12.4	12.5	
SD	1.3	6.1	12.0	1.3	5.5	10.6	10.0	13.8	14.1	1.0	4.9	9.6	12.4	12.6	
NV	0.7	3.4	11.3	0.7	2.9	10.5	10.0	13.8	13.9	1.0	4.9	9.7	12.6	12.7	
ME	0.6	5.9	13.0	0.6	5.5	11.4	10.0	14.0	14.0	1.0	4.9	9.6	13.0	13.1	
UT	0.6	4.3	11.9	0.6	3.9	10.9	10.0	14.4	14.6	1.0	4.9	9.5	13.1	13.2	
CT	0.0	4.3	12.7	0.0	4.1	11.1	10.0	14.6	14.6	1.0	4.9	9.6	13.3	13.3	
MT	1.0	6.3	12.3	1.0	5.8	11.1	10.0	14.1	14.4	1.0	4.9	9.7	13.2	13.4	
NM	0.8	5.5	12.3	0.8	4.9	11.2	10.0	14.4	14.7	1.0	4.9	9.7	13.3	13.5	
ID	1.4	5.4	12.3	1.4	5.2	11.3	10.0	14.3	14.6	1.0	4.9	9.7	13.4	13.6	
NE	0.3	5.6	12.3	0.3	5.2	11.2	10.0	14.8	15.3	1.0	4.9	9.7	13.4	13.7	
MD	0.9	4.7	12.9	0.8	4.3	11.5	10.0	14.9	15.1	1.0	4.9	9.6	13.6	13.7	
WV	1.3	6.2	13.3	1.3	5.8	12.1	10.0	14.6	14.8	1.0	5.0	9.8	13.8	14.0	
MN	1.0	6.1	13.9	1.0	5.4	12.1	10.0	15.6	16.0	1.0	4.9	9.6	13.7	14.0	
KS	0.4	6.0	13.0	0.3	5.6	11.8	10.0	15.0	15.6	1.0	5.0	9.7	13.7	14.1	
MS	1.7	6.9	13.9	1.7	6.5	12.7	10.0	15.0	15.2	1.0	5.0	9.8	14.0	14.2	
CO	0.9	5.2	12.9	0.8	4.7	11.8	10.0	15.3	15.6	1.0	4.9	9.7	14.0	14.2	
OR	1.3	5.5	12.8	1.3	5.1	11.9	10.0	15.0	15.2	1.0	5.0	9.8	14.0	14.2	
IA	0.0	6.8	13.5	0.0	6.4	12.2	10.0	15.2	15.8	1.0	5.0	9.8	13.9	14.3	
KY	0.0	6.6	13.9	0.0	6.1	12.7	10.0	15.1	15.3	1.0	4.9	9.8	14.2	14.3	
AR	0.5	6.7	13.8	0.5	6.3	12.6	10.0	15.1	15.4	1.0	5.0	9.8	14.1	14.3	
AZ	0.9	5.1	13.0	0.8	4.5	12.1	10.0	15.4	15.7	1.0	4.9	9.8	14.1	14.4	
NJ	0.2	4.7	13.0	0.2	4.1	11.6	10.0	15.8	16.0	1.0	4.9	9.7	14.4	14.5	
LA	0.3	6.1	13.5	0.3	5.5	12.4	10.0	15.3	15.5	1.0	4.9	9.7	14.3	14.5	
SC	1.0	6.2	14.1	1.0	5.8	13.1	10.0	15.3	15.5	1.0	5.0	9.8	14.4	14.5	
MA	0.0	4.7	13.1	0.0	4.3	11.6	10.0	15.5	15.7	1.0	4.9	9.7	14.4	14.6	
OK	0.0	6.2	13.6	0.0	5.7	12.3	10.0	15.4	15.8	1.0	4.9	9.8	14.2	14.6	
VA	0.7	6.0	14.0	0.7	5.6	12.8	10.0	15.5	15.7	1.0	4.9	9.7	14.4	14.6	
WA	0.1	5.8	13.6	0.1	5.4	12.5	10.0	15.5	15.8	1.0	5.0	9.8	14.4	14.6	
WI	1.1	6.6	14.1	0.9	6.0	12.5	10.0	15.9	16.2	1.0	4.9	9.8	14.4	14.6	
AL	0.9	6.7	14.4	0.8	6.3	13.1	10.0	15.6	15.8	1.0	5.0	9.8	14.5	14.7	
TN	0.2	6.5	14.5	0.3	6.1	13.3	10.0	15.7	15.9	1.0	5.0	9.8	14.7	14.9	
GA	1.0	6.5	14.6	1.0	6.1	13.5	10.0	15.6	15.9	1.0	4.9	9.8	14.6	14.9	
IN	0.7	6.5	14.3	0.7	6.1	12.8	10.0	16.0	16.4	1.0	4.9	9.7	14.7	15.1	
MO	0.6	6.5	14.1	0.6	6.1	12.9	10.0	15.8	16.3	1.0	4.9	9.8	14.7	15.1	
MI	0.3	6.4	14.6	0.3	5.9	13.2	10.0	16.4	16.7	1.0	4.9	9.8	14.9	15.2	
NY	0.3	5.3	13.5	0.3	4.1	11.8	10.0	16.3	16.6	1.0	4.8	9.7	15.0	15.2	
NC	0.9	6.9	15.0	0.9	6.6	14.0	10.0	16.1	16.3	1.0	5.0	9.8	15.2	15.4	
OH	0.6	6.6	14.7	0.6	6.2	13.4	10.0	16.4	16.8	1.0	5.0	9.8	15.2	15.6	
IL	0.8	5.7	14.0	0.8	5.1	12.6	10.0	16.6	17.2	1.0	4.9	9.7	15.4	15.8	
FL	1.0	6.6	14.6	1.0	6.2	13.7	10.0	16.5	17.0	1.0	5.0	9.8	15.5	15.9	
PA	0.8	6.4	14.6	0.9	6.0	13.2	10.0	16.7	17.2	1.0	5.0	9.8	15.5	16.0	
TX	0.9	7.3	15.2	0.9	6.7	14.1	10.0	16.7	17.6	1.0	5.0	9.8	15.4	16.0	
CA	0.8	6.6	14.6	0.8	6.2	13.9	10.0	16.9	17.6	1.0	5.0	9.9	16.1	16.7	

Table 1: Here be dragons